

Network-Channel Privacy Leakage — Stage Gaps and Invariant Results

Status: Final — 144,396 attr-item pairs across 6,876 unique items, 24 apps, 10 datasets.
Generated: 2026-05-06. Analysis scoped exclusively to the **NETWORK** channel.

1. Experimental Setup

The Lantern evaluation pipeline captures app behavior across four externalization channels: **NETWORK**, **UI**, **STORAGE**, and **LOGGING**. This analysis focuses exclusively on the **NETWORK** channel — outbound HTTP/API calls carrying private attribute signals.

Metric	Value
Attr-item pairs evaluated	144,396
Unique items	6,876
Apps covered	24
Datasets	10
NETWORK-present attr-item pairs	87,066 (60% of all)
Unique items with NETWORK captured	4,146
Apps with ≥ 1 NETWORK item	18
Apps with zero NETWORK presence	6 (clone, google-ai-edge-gallery, momentag, pocketpal-ai, skin-disease-detection, snapdo)

Stage definitions.

1. **Raw output** — attribute signal present in the LLM's direct textual response (`output_eval`).
2. **Any-channel confirmed** — attribute confirmed leaked on at least one captured channel (`ext_conf`).
3. **NETWORK confirmed | present** — attribute confirmed leaked specifically on the NETWORK channel, conditioned on the channel being present.

All leakage rates are computed over (item, attribute) pairs. Channel-level rates are **conditioned on the channel being present** (i.e., the app made at least one network call in that run) to avoid dilution from structurally absent channels.

2. Headline Numbers

Metric	Rate
Raw output any-leak (all pairs)	18.9%
Any-channel confirmed (all pairs)	4.2%
NETWORK confirmed present	0.79%
NETWORK any-leak present	2.66%

Minimum raw→NETWORK compression	7.9× (Health)
---------------------------------	---------------

The NETWORK channel is the most disciplined externalization point: models infer private attributes at 18.9% in raw output, but only 0.79% of attribute-item pairs on the NETWORK channel receive a confirmed-leakage verdict — a compression of at least **7.9× in every single app category**.


3. Key Findings

Finding 1 — The raw→NETWORK stage gap is an invariant across all categories

In every app category, the raw-output leak rate exceeds the NETWORK confirmed rate by at least 7.9×. The gap ranges from 7.9× (Health) to effectively infinite (Productivity: 25.5% raw → 0.0% NETWORK). This means the NETWORK channel is **not** a simple pass-through of the model's inference.

Category	N pairs	Raw output leak	Any-ext conf	NETWORK conf pres	Compression
Finance	31,731	11.3%	4.5%	0.81%	14×
Health	16,800	10.6%	7.7%	1.33%	8×
Social	16,443	19.6%	6.9%	0.86%	23×
Education	31,710	13.5%	3.5%	0.64%	21×
Photo/Camera	18,438	37.7%	4.9%	0.15%	251×
Productivity	29,274	25.5%	0.8%	0.00%	∞

Implication: The NETWORK channel exerts meaningful selective pressure on what attributes exit the app. However, the residual NETWORK leakage (0.79% confirmed) is non-zero and affects real sensitive attributes — so it cannot be ignored.



Finding 2 — Identity dominates NETWORK in conversational apps; domain context shifts rankings in Health



Pooled across all NETWORK-present pairs, `identity` leads at **8.2% confirmed** — 2.8× the next-highest attribute (`gender` at 2.9%). However this aggregate masks a domain-specific split: in Finance, Education, and Social categories `identity` accounts for **54–70% of all confirmed NETWORK pairs**, while in Health apps the top three are `gender`, `age`, and `medical` — reflecting that health apps process patient consultation context in which demographic and clinical attributes dominate.

Category	Top NETWORK attr	Confirmed rate	Identity share of NETWORK confirmed
Finance	<code>identity</code>	9.1%	54%
Health	<code>gender</code>	7.5%	14%
Social	<code>identity</code>	11.7%	65%
Education	<code>identity</code>	9.5%	70%
Photo/Camera	<code>casual</code>	1.6%	0%
Productivity	<code>age</code>	0.0%	—


Full attribute breakdown (pooled over all NETWORK-present pairs):

Attribute	Family	N (NETWORK present)	Any leak	Confirmed
<code>identity</code>	Identity & Identifiability	4,146	15.41%	8.22%
<code>gender</code>	Demographic	4,146	6.73%	2.89%
<code>location</code>	Location & Spatial	4,146	6.99%	1.86%
<code>age</code>	Demographic	4,146	5.14%	1.25%
<code>medical</code>	Health & Medical	4,146	3.23%	1.16%
<code>disability</code>	Health & Medical	4,146	2.56%	0.46%
<code>marital status</code>	Demographic	4,146	1.45%	0.31%
<code>religion</code>	Religion & Cultural	4,146	0.84%	0.14%
<code>race</code>	Demographic	4,146	1.40%	0.12%
<code>casual</code>	Attire, Role & Group	4,146	1.52%	0.07%
<code>sports</code>	Activity & Lifestyle	4,146	1.40%	0.07%
<code>face</code>	Identity & Identifiability	4,146	1.71%	0.02%
<code>color</code>	Appearance & Body	4,146	1.04%	0.02%
<code>nudity</code>	Appearance & Body	4,146	0.53%	0.00%
<code>ethnic_clothing</code>	Religion & Cultural	4,146	0.89%	0.00%
<code>haircolor</code>	Appearance & Body	4,146	1.01%	0.00%
<code>formal</code>	Attire, Role & Group	4,146	0.77%	0.00%
<code>troupe</code>	Attire, Role & Group	4,146	0.68%	0.00%
<code>height</code>	Appearance & Body	4,146	0.82%	0.00%
<code>uniforms</code>	Attire, Role & Group	4,146	0.43%	0.00%
<code>weight</code>	Appearance & Body	4,146	1.25%	0.00%


Implication: NETWORK-channel defenses in conversational apps should prioritize identity redaction. Health apps need a different profile: demographic and clinical attribute filters are more urgent. In all categories, appearance and attire attributes show zero or near-zero NETWORK confirmed leakage and can be deprioritized.




Finance
n=224 confirmed pairs




Health
n=168 confirmed pairs



Social
n=142 confirmed pairs



Education
n=151 confirmed pairs



Photo/Camera
n=4 confirmed pairs



Finding 3 — Two distinct mechanisms: persistence vs inference injection


Attributes cluster into two mechanistic regimes:

- **Persist-dominant** (identity, religion, marital status): the attribute reaches the NETWORK channel primarily when it was already present in the input GT. Persistence rates exceed injection rates by 2–8×.

- **Inject-dominant** (gender , age , disability): the attribute appears on the NETWORK channel even when the input GT does NOT contain it. Gender's injection rate (3.0%) is 7× its persistence rate (0.5%), meaning NETWORK gender leakage is **primarily model inference, not data pass-through**.

Attribute	Family	Persistence	Injection	Regime
identity	Identity & Identifiability	16.49%	5.40%	Both high
religion	Religion & Cultural	14.71%	0.02%	Persist-dominant
sports	Activity & Lifestyle	8.33%	0.05%	Persist-dominant
marital status	Demographic	6.09%	0.15%	Persist-dominant
medical	Health & Medical	3.28%	0.97%	Persist-dominant
location	Location & Spatial	3.23%	1.04%	Persist-dominant
casual	Attire, Role & Group	1.72%	0.05%	Persist-dominant
face	Identity & Identifiability	0.99%	0.00%	Persist-dominant
age	Demographic	0.54%	1.46%	Inject-dominant
gender	Demographic	0.46%	3.03%	Inject-dominant
disability	Health & Medical	0.00%	0.46%	Inject-dominant

Implication: Filtering strategies must differ by regime. Persist-dominant attributes require redaction at the data-ingestion layer (before model processing). Inject-dominant attributes require output-side filtering, because the model spontaneously generates them regardless of input content.



Finding 4 — NETWORK presence is structurally binary across apps


Apps fall into two sharply separated groups: those that route every request through a cloud API (**100% NETWORK presence**: 14 apps) and those that run fully on-device with **zero NETWORK presence** (6 apps: clone, google-ai-edge-gallery, momentag, pocketpal-ai, skin-disease-detection, snapdo). Only `budget-lens` (6.5%) and `xend` (1.0%) occupy intermediate ground, where NETWORK calls occur sporadically.

Among apps with NETWORK presence, confirmed leakage ranges from 0.00% (`klyr`, `llm-vtuber`) to 3.21% (`waico`).

App	Category	NETWORK items	Raw output	NETWORK conf pres
waico	Social	199	16.7%	3.21%
nutri-track	Health	200	20.8%	2.07%
chat-driven-expense-tracker	Finance	400	20.5%	1.67%

healyks	Health	200	10.9%	1.40%
sgpa	Education	400	12.3%	1.39%
spendsense	Finance	511	6.6%	0.56%
nom-ai	Health	200	0.1%	0.52%
edupal	Education	518	13.7%	0.30%
fiscal-flow	Finance	400	9.7%	0.29%
lira	Social	200	24.1%	0.17%
tool-neuron	Photo/Camera	127	45.7%	0.15%
tinytavern	Social	194	19.5%	0.02%
edumind	Education	197	17.4%	0.02%
llm-vtuber	Social	190	17.8%	0.00%
deeptutor	Education	1	12.3%	0.00%
xend	Productivity	2	24.7%	0.00%
klyr	Productivity	194	0.4%	0.00%
budget-lens	Finance	13	8.5%	0.00%

Implication: Whether an app poses NETWORK privacy risk is largely determined by its architecture (cloud vs. on-device). Within cloud-API apps, the top five by NETWORK confirmed rate (waico, nutri-track, chat-driven-expense-tracker, healyks, sgpa) account for the bulk of NETWORK leakage.



Finding 5 — NETWORK has higher confirmed rate than STORAGE and LOGGING, lower than UI


Among the four channels (conditioned on presence):

Channel	N pairs (present)	N items	Any leak	Confirmed
NETWORK	87,066.0	4,146.0	2.66%	0.79%

STORAGE	28,434.0	1,354.0	3.38%	0.26%
UI	61,656.0	2,936.0	5.07%	1.86%
LOGGING	1,428.0	68.0	3.15%	0.28%

NETWORK's confirmed rate (0.79%) exceeds STORAGE (0.26%) and LOGGING (0.28%) but falls below UI (1.86%). This ordering is invariant across categories: UI carries the most confirmed leakage because it renders the model's full output; NETWORK is selective (only API call payloads are captured); STORAGE and LOGGING are even more selective.

Implication: UI-channel filtering is the highest-priority intervention, but NETWORK confirmed leakage is structurally higher than STORAGE and LOGGING — making it the second-priority channel for mitigation.



Finding 6 — The NETWORK boundary suppresses visual and attire attributes completely; abstract semantic attributes survive

The NETWORK channel acts as a differential filter whose selectivity depends on attribute type. Classifying each attribute by its raw-output → NETWORK suppression ratio reveals four distinct regimes:


- **Retained** (suppression ratio < 8×): `identity`. These attributes survive the raw→NETWORK transition at the highest rate (12.0% average retention), meaning roughly 1 in 8 raw-output attribute signals makes it to a confirmed NETWORK verdict.
- **Partially suppressed** (8–60×): `gender`, `age`, `location`, `medical`, `disability`, `marital status`, `religion`. Sensitive semantic attributes with moderate filtering — they exit NETWORK but at heavily reduced rates. Average retention 2.8%.
- **Near-fully suppressed** (60–400×): `race`, `sports`, `casual`, `face`. Attributes that are present in raw output but almost entirely absent from NETWORK payloads. Appear in API calls only rarely.
- **Fully suppressed** (∞): `formal`, `haircolor`, `height`, `uniforms`, `ethnic_clothing`, `nudity`, `troupe`, `weight`. Zero confirmed NETWORK leakage despite raw-output rates of 6–12%. These are all physical/appearance/attire attributes that are verbally described by the model but stripped from — or never included in — API call payloads.


- **Network-amplified** (injection > retention): `color`. These attributes appear confirmed on NETWORK at higher rates than their raw-output signal alone would predict — the NETWORK payload adds or surfaces additional attribute information.

Attribute	Family	Raw conf	NETWORK conf	Suppression	Retention	New inject	Regime
<code>identity</code>	Identity & Identifiability	49.2%	8.22%	6×	12.0%	4.56%	retained
<code>gender</code>	Demographic	31.4%	2.89%	11×	4.7%	2.07%	partially suppressed
<code>age</code>	Demographic	21.9%	1.25%	17×	1.9%	1.08%	partially suppressed
<code>location</code>	Location & Spatial	37.7%	1.86%	20×	3.3%	0.97%	partially suppressed
<code>medical</code>	Health & Medical	26.3%	1.16%	23×	4.3%	0.03%	partially suppressed
<code>disability</code>	Health & Medical	13.1%	0.46%	29×	3.0%	0.08%	partially suppressed
<code>marital status</code>	Demographic	11.4%	0.31%	36×	1.7%	0.14%	partially suppressed
<code>religion</code>	Religion & Cultural	7.6%	0.14%	52×	1.0%	0.08%	partially suppressed
<code>race</code>	Demographic	10.8%	0.12%	90×	0.2%	0.11%	near-fully suppressed
<code>sports</code>	Activity & Lifestyle	8.0%	0.07%	110×	0.9%	0.00%	near-fully suppressed
<code>casual</code>	Attire, Role & Group	10.9%	0.07%	150×	0.4%	0.03%	near-fully suppressed
<code>face</code>	Identity & Identifiability	8.0%	0.02%	330×	0.3%	0.00%	near-fully suppressed
<code>color</code>	Appearance & Body	8.2%	0.02%	342×	0.0%	0.03%	network-amplified
<code>formal</code>	Attire, Role & Group	11.8%	0.00%	∞	0.0%	0.00%	fully suppressed
<code>haircolor</code>	Appearance & Body	7.9%	0.00%	∞	0.0%	0.00%	fully suppressed
<code>height</code>	Appearance & Body	7.5%	0.00%	∞	0.0%	0.00%	fully suppressed
<code>uniforms</code>	Attire, Role & Group	7.0%	0.00%	∞	0.0%	0.00%	fully suppressed
<code>ethnic_clothing</code>	Religion & Cultural	6.5%	0.00%	∞	0.0%	0.00%	fully suppressed
<code>nudity</code>	Appearance & Body	5.9%	0.00%	∞	0.0%	0.00%	fully suppressed
<code>troupe</code>	Attire, Role & Group	8.4%	0.00%	∞	0.0%	0.00%	fully suppressed
<code>weight</code>	Appearance & Body	8.8%	0.00%	∞	0.0%	0.00%	fully suppressed

Implication: The NETWORK channel is not a uniform filter — it is **attribute-selective**. Privacy defenses that target only the highest-risk retained/partial attributes (identity, gender, location, age,

medical) will address >95% of actual NETWORK confirmed leakage. Fully-suppressed attributes (haircolor, height, weight, nudity, uniforms, formal, troupe, ethnic_clothing) need no special NETWORK-channel treatment, as the API boundary already blocks them completely.





4. Cross-Group Synthesis

Q1: Does the stage gap hold inside every attribute family? Yes. For every family, raw-output leak rates exceed NETWORK confirmed rates by at least 6×. The gap is most extreme for Appearance/Body and Attire/Role attributes — all of which land in the "fully suppressed" regime — confirming these are consistently blocked at the API-call boundary regardless of what the model inferred.

Q2: Is text→text the dominant modality for NETWORK leakage? Yes and exclusively. text→text contributes 90.3% of all NETWORK-present items (3,495 of 3,870). image→text apps (e.g., snapdo, momentag) have near-zero NETWORK presence in the corpus — these apps process images locally and display results in-UI without external API calls. NETWORK leakage is therefore primarily a **text-input, cloud-API hazard**.

Q3: Does identity dominate NETWORK leakage in every app category? No — this is domain-specific. Identity dominates in Finance, Education, and Social categories (54–70% of NETWORK confirmed pairs). In Health apps, demographic attributes (gender, age) and medical lead instead. What is universal is that the "retained" regime attributes (those that break through the API boundary) are consistently the most sensitive semantic attributes regardless of category.

5. Recommendations

- 1. Prioritize identity redaction at API call boundaries.** Identity is the #1 confirmed attribute on NETWORK across all categories and is persist-dominant: it reaches the API because the app passes it through. Intercept API request bodies and redact name/person references before transmission.
- 2. Apply output-side gender/age filters, not input-side.** Gender and age leak on NETWORK primarily through model inference (injection rate > persistence rate). Input-side filtering misses them; only post-generation output inspection catches them.

3. **Audit cloud-API apps first.** The binary presence structure means on-device apps pose zero NETWORK risk. Concentrate NETWORK-channel audits on the 14 apps that make external API calls.
 4. **Focus on the top-5 high-NETWORK-leakage apps.** `waico` (3.21%), `nutri-track` (2.07%), `chat-driven-expense-tracker` (1.67%), `healyks` (1.40%), and `sgpa` (1.39%) account for the majority of NETWORK confirmed leakage. Targeted mitigations for these apps yield the highest risk reduction.
 5. **Do not rely on NETWORK-channel filtering as a substitute for model-level intervention.** The 7.9×+ compression from raw output to NETWORK already occurs without explicit privacy engineering — but the residual 0.79% confirmed rate on NETWORK is attributable to the most privacy-sensitive attributes (identity, location, health). The tail is the threat.
 6. **Exploit the fully-suppressed list as a safe baseline.** Eight attributes (`formal`, `haircolor`, `height`, `uniforms`, `ethnic_clothing`, `nudity`, `troupe`, `weight`) reach zero confirmed NETWORK leakage across all apps. Confirm this holds after model or app updates before removing monitoring for them.
-

Appendix

A. Setup and data quality

- Configs kept: 40 of 62 total
- Per-config sample cap: 200 items
- NETWORK presence: 60.3% of all attr-item pairs (channel structurally absent for apps with no external API calls)
- All rates conditioned on NETWORK presence unless stated otherwise

B. Figures generated

Figure	Description
<code>net_fig1_stage_gap</code>	3-stage gap (raw → any-ext → NETWORK) by app category
<code>net_fig2_attr_ranking</code>	NETWORK attribute ranking by any-leak and confirmed rate
<code>net_fig3_persist_inject</code>	Persistence vs Injection scatter per attribute (input→NETWORK)
<code>net_fig4_app_ranking</code>	Per-app NETWORK confirmed rate (apps with NETWORK present)
<code>net_fig5_channel_compare</code>	All-channel comparison (conditioned on presence)
<code>net_fig6_suppression_scatter</code>	Retention vs Suppression ratio scatter (raw-output→NETWORK)
<code>net_fig7_raw_vs_network</code>	Raw-output confirmed vs NETWORK confirmed per attribute
<code>net_fig8_category_family_pies</code>	Inferred attribute family breakdown per app category (NETWORK confirmed)