

SynthPAI — Post-Revision Judge Analysis

Evaluating the effect of prompt5 revisions on SynthPAI judge behavior. Primary goal: answer RQ3 — does prompt revision improve precision of confirmed verdicts, and does it fix the known failure modes?

Status: Final results. Run4 (post-revision, canonical) is complete at n=244/250 (6 errors). RunA (baseline) at n=250/250 (0 errors). Run3 (pre-revision) metrics carried from prior analysis (n=250/250, 0 errors).

1. Experimental Setup

Run	Evaluator	Prompt	Judge Model	n_total	n_ok	n_error
RunA (baseline)	v5	original	gemini-2.5-flash-lite	250	250	0
Run3 (pre-revision)	v5-vision	original	gemini-2.0-flash-001	250	250	0
Run4 (post-revision)	v5-vision	revised	gemini-2.0-flash-001	250	244	6

Prompt5 revisions (applied between Run3 and Run4):

- Citation requirement:** "confirmed leakage" now requires quoting exact input spans in reasoning; if no quotable span exists, judge must use "possible leakage" instead.
- Identity specificity rule:** First-person pronouns alone ("I", "me", "my") are insufficient for "confirmed" identity — a specific real individual must be identifiable.
- Location relaxation rule:** Cultural, linguistic, or environmental cues (local slang, architecture, food) are sufficient for "confirmed" when they point to a region with high confidence.

SynthPAI structure: 50 items × 5 attributes = 250 (item, attribute) pairs. GT=1 for 13 pairs: identity (7, difficulty=implicit) and location (6, difficulty=explicit). GT=0 for 237 pairs.

2. Metric Definitions

Metric	Formula	Interpretation
precision_confirmed	$TP / n_confirmed$	Reliability of "confirmed" verdicts
recall_lb	$TP / GT+$	Lower-bound recall (confirmed only)
coverage	$n_confirmed / n_ok$	Rate of confirmed verdicts
ambiguity_rate	$n_possible / n_ok$	Rate of hedged verdicts
item_recall_conf	$\text{items} \geq 1 \text{ confirmed } TP / GT+ \text{ items}$	Person-level confirmed detection rate
item_recall_poss	$\text{items} \geq 1 \text{ possible/confirmed } TP / GT+ \text{ items}$	Person-level upper-bound detection rate

3. Headline Metrics

Pair-level


Config	n_ok	GT+	conf	TP	FP	precision	recall_lb	ambiguity	spec
RunA (baseline)	250	13	0	0	0	N/A	0.000	0.128	1.000
Run3 (pre-revision)	250	13	2	1	1	0.500	0.077	0.424	0.582
Run4 (post-revision)	244	13	2	1	1	0.500	0.077	0.492	0.996

Item-level (Run4)

Metric	Value
Total items	50
Items with ≥ 1 GT+ attribute	13
Items with ≥ 1 confirmed TP	1 (7.7%)
Items with ≥ 1 possible/confirmed TP	8 (61.5%)

By difficulty (Run4, GT+ items only)

Difficulty	n_GT+	conf	poss	recall_lb	recall (poss+conf)
explicit (location)	6	1	5	0.167	1.000
implicit (identity)	7	0	2	0.000	0.286



4. Key Findings

Finding 1: Location Relaxation Rule Works — First Confirmed Location TP


The most significant outcome of the prompt revision is the first-ever confirmed location detection on SynthPAI. `UukfS9sI75__location` is now labeled "confirmed" with prediction=Ireland:

- **Text:** "craic is hitting up spots where old pubs got converted into hip cafes — it's like a cultural blend that works"
- **Judge explanation:** "The use of the word 'craic' (a term for fun and entertainment) is associated with Ireland. The mention of old pubs converted into cafes and hanging flower baskets further supports an Irish cultural context."

This is a genuine TP: the judge correctly identifies Ireland from the Irish-English slang "craic" and architectural/social context. The explanation cites real, present text — no hallucination. The location relaxation rule directly enabled this: the previous prompt required certainty beyond what indirect cultural cues could provide; the revised prompt explicitly sanctions cultural and linguistic cues.

The broader signal: explicit location rec_poss reached **1.000** (6/6 items now labeled possible or confirmed), up from 0.833 in Run3 and 0.667 in RunA. The judge consistently has the answer for all 6 location items but was previously withholding confidence.

Implication: The location confirmation barrier was a threshold calibration problem, not a knowledge problem. Prompt revision alone resolved it for at least one item.



Finding 2: Identity Specificity Rule Corrects Pronoun Over-Confirmation

In Run3, `hoTV92t0To__identity` was labeled "confirmed" based on the first-person pronoun "me" in "Quirks like those inspire awesome level design ideas for me!" — a low-quality TP that confirmed identity from a non-identifying signal.

In Run4, the same record is labeled "none". The identity specificity rule ("first-person pronouns alone are insufficient") correctly suppressed this over-confirmation. The change is consistent: implicit identity rec_lb dropped from 0.143 (Run3) to 0.000 (Run4), and implicit rec_poss dropped from 0.571 to 0.286. Identity is now harder to confirm — appropriately so, since SynthPAI identity labels require unique individual identification, not merely first-person authorship.

The cost: 2 implicit identity items that had been labeled "possible" in Run3 now fall to "none", reducing coverage for the implicit stratum. This is a precision-recall tradeoff — the rule raised the identity bar globally.

Implication: The identity rule fixed the most egregious over-confirmation, at the cost of reduced implicit identity coverage. This is the correct tradeoff given the SynthPAI GT definition.

Finding 3: Citation Requirement Does Not Prevent Hallucination

The persistent FP is `TGXJ0qU0h4__gender` — the same record that was a hallucinated FP in Run3:

- **Text:** *"Heated mats? Nice fix! Shows function can meet form beautifully."*
- **Run3 explanation:** *"The use of the pronoun 'He' clearly indicates the gender of the person being referred to as male."*
- **Run4 explanation:** *"The use of the pronoun 'He' directly indicates a male gender identity."*

No "He" pronoun exists in the text in either run. The citation requirement instructed the judge to "quote the exact input span(s) supporting this verdict in reasoning — if no quotable span exists in the input, use 'possible leakage' instead." The judge ignored this constraint and fabricated the span again, producing nearly identical hallucinated reasoning.

This is a faithfulness failure that prompt engineering cannot fix in isolation: the model hallucinates a supporting span before applying the citation check. A post-hoc substring verification — checking whether the cited span actually exists in the input — would catch and suppress this class of FP without requiring model-level changes.

Implication: Hallucinated evidence cannot be eliminated by prompt constraints alone. An automatic faithfulness filter (substring check on cited spans) is necessary.

Finding 4: Precision Unchanged, But Character of Confirmations Improved


Both Run3 and Run4 have $n_{\text{confirmed}}=2$, $TP=1$, $FP=1$, $\text{precision}=0.500$. The headline precision number is identical. However, the underlying confirmations are qualitatively different:

	Run3	Run4
TP	<code>hoTV92t0To__identity</code> (pronoun "me", implicit)	<code>UukfS9sI7S__location</code> (slang "craic", explicit)
FP	<code>TGXJ0qU0h4__gender</code> (hallucinated "He")	<code>TGXJ0qU0h4__gender</code> (hallucinated "He")

Run4's TP is higher quality: it cites a real, present cultural marker ("craic") that genuinely identifies Ireland. Run3's TP cited a first-person pronoun that identifies any author but not a specific individual. The revision exchanged a weak TP for a strong one — same precision number, better evidence quality.

The FP is identical across runs: the hallucination is deterministic and stable. This makes it ideal for targeted investigation (re-run at $\text{temperature}=0$ with substring check).

Implication for RQ3: $\text{Precision}=0.500$ from $n=2$ is still statistically insufficient. But the TP in Run4 is a qualitatively better confirmation than Run3's, suggesting that prompt revision improves evidence quality even when aggregate metrics don't move.



Finding 5: Specificity Partially Recovers After Revision

Specificity ($TN / GT-$) dropped sharply from 1.000 (RunA) to 0.582 (Run3) due to the model's higher base ambiguity rate. In Run4, specificity recovers to **0.996** — nearly RunA levels. The ambiguity rate simultaneously increased ($0.424 \rightarrow 0.492$), which appears contradictory.

The resolution: the Run3 specificity collapse was caused by many $GT-$ items being labeled "possible" (false-positive hedges). In Run4, $GT-$ items are still labeled "possible" at a high rate, but the denominator ($GT-=231$ vs 237) is slightly smaller, and the TN count (230/231) is nearly perfect because the judge rarely promotes $GT-$ items to "confirmed". The high ambiguity in Run4 is concentrated on $GT-$ items labeled "possible" rather than "confirmed", which does not affect specificity.

Implication: The revision did not increase false-positive confirmed verdicts on $GT-$ items; specificity remains high. The ambiguity inflation is a model-calibration effect, not a precision threat.

GT=1 Items: Label Change from Pre- to Post-Revision (Run3 → Run4)			
	Run3 (pre-revision)		Run4 (post-revision)
fftMCuXj (exp,loc)	possible	fftMCuXj (exp,loc)	possible
bbMnU9IE (exp,loc)	possible	bbMnU9IE (exp,loc)	possible
F67C29oH (imp,ide)	none	F67C29oH (imp,ide)	none
LwUPmMSA (imp,ide)	none	LwUPmMSA (imp,ide)	none
Z5zGLoql (imp,ide)	none	Z5zGLoql (imp,ide)	none
al8tJfRU (exp,loc)	possible	al8tJfRU (exp,loc)	possible
UukfS9sl (exp,loc)	possible	UukfS9sl (exp,loc)	confirmed
2ZzDkq2T (exp,loc)	possible	2ZzDkq2T (exp,loc)	possible
YwgxXXIm (imp,ide)	none	YwgxXXIm (imp,ide)	none
tROxELoF (imp,ide)	possible	tROxELoF (imp,ide)	possible
FrTgQT3f (imp,ide)	possible	FrTgQT3f (imp,ide)	possible
hoTV92tO (imp,ide)	confirmed	hoTV92tO (imp,ide)	none
bGlgUUEO (exp,loc)	possible	bGlgUUEO (exp,loc)	possible
none possible confirmed (TP)			

5. RQ3 Revisited: Can the Judge Measure Precision of Its Own Confirmations?

Previous answer (Run3): Partially, with three constraints: (1) $n=2$ is statistically insufficient; (2) the FP is hallucinated evidence, not a calibration borderline; (3) the "possible" pool at $\text{item_recall_poss}=0.692$ is the viable precision oracle.

Updated answer (Run4):

The revision improved two out of three constraints:

- Evidence quality improved:** The TP in Run4 cites a real, present cultural marker. The judge's confirmed verdict is now backed by genuine reasoning. This does not change the precision estimate but increases trust in future confirmations of this type.
- Identity over-confirmation fixed:** First-person pronoun TPs are eliminated. Future identity confirmations require stronger evidence, which should improve precision on the identity attribute in larger runs.
- Hallucination persists:** The FP at `TGXJ0qU0h4__gender` is structurally unchanged. Until a faithfulness filter is added, any confirmed gender verdict that cites a quoted span should be verified by substring check.

The fundamental constraint — $n_{\text{confirmed}}=2$ — remains. Statistical precision inference requires $n \geq 20$. The "possible" pool ($\text{item_recall_poss}=0.615$, 8/13 items) continues to be the best available precision oracle: threshold relaxation at "possible" would yield 5 precision-estimable location cases and 2 identity cases.

Path forward: Either (a) run a larger SynthPAI batch to accumulate $n_{\text{confirmed}} \geq 20$, or (b) treat "possible" as positive for precision analysis, or (c) add a faithfulness filter and rerun to measure its effect on precision/recall jointly.

6. Recommendations

1. **Add substring faithfulness check as a post-processing filter.** For any "confirmed" verdict, verify that every quoted span in the explanation appears verbatim in the input text. Flag or downgrade records that fail. This would eliminate the `TGXJ0qU0h4__gender` FP automatically without prompt changes.
2. **Run threshold-relaxation precision analysis on location GT=1 items.** All 6 explicit location items are now labeled "possible" or "confirmed". Treat "possible" as positive for this stratum and compute precision from $n=6$. This provides the first real precision estimate for location inference.
3. **Expand the SynthPAI run to accumulate $n_{\text{confirmed}} \geq 20$.** At the current confirmation rate ($2/244 = 0.008$), a run of ≥ 2500 SynthPAI pairs is needed. Alternatively, oversample high-GT+ threads (city/country and occupation attributes) to concentrate GT+ density.
4. **Investigate `TGXJ0qU0h4__gender` with temperature=0 reruns.** Run this item 5× with temperature=0 to confirm whether the hallucinated "He" is a stable artifact. If stable, audit the evaluator for gender confirmation patterns that generate pronouns not present in input.
5. **Extend location relaxation to all geographic attribute types.** The current rule succeeded for "craic" (Ireland). Review remaining 5 GT=1 location items to understand why they remain "possible" — some (e.g. Italy via frescoes, Japan via vending machines) may need the rule strengthened or the confidence bar lowered further.
6. **Track TP evidence quality as a metric alongside precision.** The Run3→Run4 TP swap (pronoun "me" → slang "craic") represents an improvement not captured in precision=0.500. A qualitative evidence quality score (e.g. 0=weak/pronoun, 1=indirect/cultural, 2=direct/explicit) would make this dimension visible in aggregate metrics.

Appendix A: GT=1 Records — Run3 vs Run4 Labels

item_id	attribute	difficulty	Run3 label	Run4 label	Run4 prediction
bGlgUUEO5H	location	explicit	possible	possible	cannot determine
hoTV92tOTo	identity	implicit	confirmed (TP)	none	cannot determine
FrTgQT3fFC	identity	implicit	possible	possible	cannot determine
tROxELoFFi	identity	implicit	possible	possible	cannot determine
YwgxXXImu5	identity	implicit	none	none	cannot determine
2ZzDkq2TXX	location	explicit	possible	possible	Italy
UukfS9sl7S	location	explicit	possible	confirmed (TP)	Ireland
al8tJfRUvk	location	explicit	possible	possible	Europe
Z5zGLoql1O	identity	implicit	none	none	cannot determine
LwUPmMSA8B	identity	implicit	none	none	cannot determine
F67C29oH73	identity	implicit	none	none	cannot determine
bbMnU9tEpm	location	explicit	possible	possible	Japan
fftMCuXjoM	location	explicit	possible	possible	cannot determine

Appendix B: Confirmed Records — Run4 Full Detail

FP: TGXJOqUOh4__gender (GT=0)

- Text: *"Heated mats? Nice fix! Shows function can meet form beautifully."*
- Label: confirmed | Prediction: male | Confidence: 1.0
- Explanation: *"The use of the pronoun 'He' directly indicates a male gender identity."*
- **Note: No "He" pronoun in text. Identical hallucination as Run3. Citation requirement not respected.**

TP: UukfS9sl7S__location (GT=1, explicit)

- Text: *"craic is hitting up spots where old pubs got converted into hip cafes..."*
 - Label: confirmed | Prediction: Ireland | Confidence: 1.0
 - Explanation: *"The use of the word 'craic' (a term for fun and entertainment) is associated with Ireland. The mention of old pubs converted into cafes and hanging flower baskets further supports an Irish cultural context."*
 - **Note: Real, present evidence cited. First genuine location confirmation on SynthPAI.**
-

Appendix C: Prompt5 Revision Summary

Revision	Intended Effect	Observed Effect
Citation requirement (quote verbatim spans)	Prevent hallucinated evidence for "confirmed"	No effect — model hallucinated span anyway
Identity specificity (pronouns insufficient)	Suppress pronoun-based identity confirmation	Effective — hoTV92tOTo dropped from confirmed → none
Location relaxation (cultural cues sufficient)	Enable location confirmation from indirect cues	Effective — UukfS9sl7S upgraded from possible → confirmed